

# Assignment 3.1 [Python & R]

## University of San Diego

### ADS 502

### Dingyi Duan

For Exercises 28–34, work with the churn data set.

28. Partition the data set, so that 67% of the records are included in the training data set and 33% are included in the test data set. Use a bar graph to confirm your proportions.

```
In [1]: import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: %load_ext rpy2.ipython
```

#### Python Packages

```
In [3]: import random
import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt

from sklearn.tree import DecisionTreeClassifier, export_graphviz
from statsmodels.stats.proportion import proportions_ztest
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import confusion_matrix
```

#### R Package

```
In [4]: %%R
library(readr)
library(ggplot2)
library(C50)
library(e1071)
```

R[write to console]: RStudio Community is a great place to get help:  
<https://community.rstudio.com/c/tidyverse> (<https://community.rstudio.com/c/tidyverse>)

## Python

```
In [5]: churn_py = pd.read_csv("D:/2021-Spring-textbooks/ADS-502/Website Data Sets/churn_
```

```
In [6]: churn_py.head()
```

Out[6]:

	State	Account Length	Area Code	Phone	Intl Plan	VMail Plan	VMail Message	Day Mins	Day Calls	Day Charge	...	Eve Charge	Night Mins
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	16.78	244.7
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	16.62	254.4
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	10.30	162.6
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...	5.26	196.9
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...	12.61	186.9

5 rows × 22 columns



```
In [7]: churn_train_py, churn_test_py = train_test_split(churn_py, test_size = 0.33, rand
```

```
In [8]: print ('churn_train_py has %s records'% churn_train_py.shape[0],
            '\nchurn_test_py has %s records'% churn_test_py.shape[0])
```

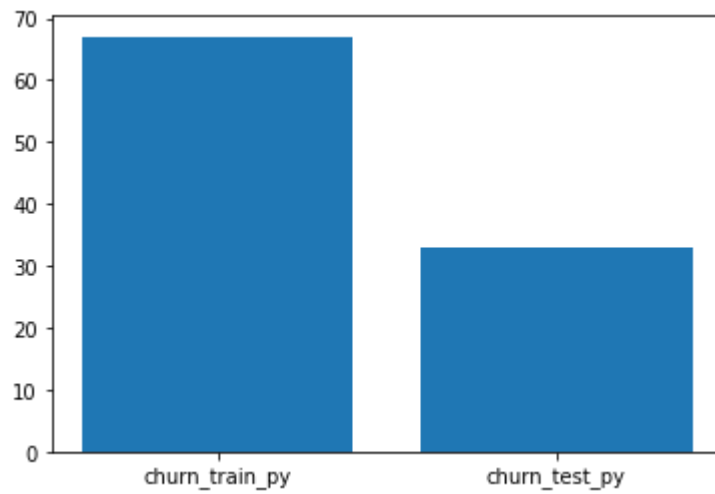
churn\_train\_py has 2233 records  
 churn\_test\_py has 1100 records

```
In [9]: print ('The proportion for churn_train_py is %s'% round(churn_train_py.shape[0]*1
            '\nThe proportion for churn_test_py is %s'% round(churn_test_py.shape[0]*1
```

The proportion for churn\_train\_py is 67.0  
 The proportion for churn\_test\_py is 33.0

```
In [10]: names_py = ['churn_train_py', 'churn_test_py']  
values_pct_py = [round(churn_train_py.shape[0]*100/churn_py.shape[0],0),round(churn_test_py.shape[0]*100/churn_py.shape[0],0)]  
plt.bar(names_py, values_pct_py)
```

Out[10]: <BarContainer object of 2 artists>



**R**

```
In [11]: %%R
churn_r <- read.csv(file = "D:/2021-Spring-textbooks/ADS-502/Website Data Sets/churn_r.csv")
set.seed(7)
n <- dim(churn_r)[1]
head(churn_r)
```

	State	Account.Length	Area.Code	Phone	Intl.Plan	VMail.Plan	VMail.Message
1	KS	128	415	382-4657	no	yes	25
2	OH	107	415	371-7191	no	yes	26
3	NJ	137	415	358-1921	no	no	0
4	OH	84	408	375-9999	yes	no	0
5	OK	75	415	330-6626	yes	no	0
6	AL	118	510	391-8027	yes	no	0

	Day.Mins	Day.Calls	Day.Charge	Eve.Mins	Eve.Calls	Eve.Charge	Night.Mins
1	265.1	110	45.07	197.4	99	16.78	244.7
2	161.6	123	27.47	195.5	103	16.62	254.4
3	243.4	114	41.38	121.2	110	10.30	162.6
4	299.4	71	50.90	61.9	88	5.26	196.9
5	166.7	113	28.34	148.3	122	12.61	186.9
6	223.4	98	37.98	220.6	101	18.75	203.9

	Night.Calls	Night.Charge	Intl.Mins	Intl.Calls	Intl.Charge	CustServ.Calls
1	91	11.01	10.0	3	2.70	1
2	103	11.45	13.7	3	3.70	1
3	104	7.32	12.2	5	3.29	0
4	89	8.86	6.6	7	1.78	2
5	121	8.41	10.1	3	2.73	3
6	118	9.18	6.3	6	1.70	0

	Old.Churn	Churn
1	False	False
2	False	False
3	False	False
4	False	False
5	False	False
6	False	False

```
In [12]: %%R
train_ind <- runif(dim(churn_r)[1]) <= 0.67 # Determin training data proportion
```

```
In [13]: %%R
# Split training and test datasets
churn_train_r <- churn_r[ train_ind, ]
churn_test_r <- churn_r[ !train_ind, ]
```


```
In [14]: %%R
dim(churn_train_r)[1]

[1] 2227
```

```
In [15]: %%R
dim(churn_test_r)[1]

[1] 1106
```

```
In [16]: %%R
values_pct_r <- c((dim(churn_train_r)[1])*100/dim(churn_r)[1],(dim(churn_test_r)[1])*100/dim(churn_r)[1])
names_r <- c("churn_train_r","churn_test_r")
barplot(values_pct_r,names.arg=names_r,col = rainbow(2))
```



29. Identify the total number of records in the training data set and how many records in the training data set have a churn value of true.

## Python

```
In [17]: print('Total number of records in the training data set is ',churn_train_py.shape[0])
Total number of records in the training data set is 2233
```

```
In [18]: print('Number of records in the training data set have a churn value of true is ',churn_train_py[churn_train_py['churn']==True].shape[0])
Number of records in the training data set have a churn value of true is 320
```

```
In [19]: # Or
churn_train_py['Churn'].value_counts()
```

```
Out[19]: False    1913
         True      320
         Name: Churn, dtype: int64
```

## R

```
In [20]: %%R
n_train_r <- dim(churn_train_r)[1]
print(paste('Total number of records in the training data set is ',n_train_r))

[1] "Total number of records in the training data set is 2227"
```

```
In [21]: %%R
a <- table(churn_train_r$Churn)["True"]
print(paste('Number of records in the training data set have a churn value of true is ',a))

[1] "Number of records in the training data set have a churn value of true is 325"
```

```
In [22]: %%R
# Or
table(churn_train_r$Churn)
```

```
False  True
 1902   325
```

**30. Use your answers from the previous exercise to calculate how many true churn records you need to resample in order to have 20% of the rebalanced data set have true churn values.**

## Python

```
In [23]: # Ratio of "True" to all Churn in churn_train
churn_train_py['Churn'].sum()*100 / churn_train_py.shape[0]
```

```
Out[23]: 14.33049708911778
```

Total churn\_train has 2233, Churn == 'True' has 320, which is 14.3% of all samples; We need 20%.

```
In [24]: print('We need %s records with Churn == True to add to our training set' % int((20 - 14.33) * 2233))

We need 158 records with Churn == True to add to our training set
```

## R

```
In [25]: %%R

print(paste('Current ratio of "Churn == True" to all training set is ', a*100/n_t

[1] "Current ratio of \"Churn == True\" to all training set is  14.593623709025
6"
```

```
In [26]: %%R

# We need 20%
print(paste('The number of Churn == True we still need:', floor((0.2*(2233)-320),

[1] "The number of Churn == True we still need: 158"
```

**31. Perform the rebalancing described in the previous exercise and confirm that 20% of the records in the rebalanced data set have true churn values.**

## Python

```
In [27]: to_resample_py = churn_train_py.loc[churn_train_py['Churn'] == True]
```

```
In [28]: # Sample from our records of interest
```

```
In [29]: our_resample_py = to_resample_py.sample(n = 158, replace = True)
```

```
In [30]: # Add the resampled records to our original training data set
```

```
In [31]: churn_train_rebal_py = pd.concat([churn_train_py, our_resample_py], axis=0)
```

```
In [32]: churn_train_rebal_py['Churn'].sum() #The current Churn == True records
```

```
Out[32]: 478
```

```
In [33]: # Or
churn_train_rebal_py['Churn'].value_counts()
```

```
Out[33]: False    1913
         True      478
         Name: Churn, dtype: int64
```

```
In [34]: # Proportion of our desired records in the new (rebalanced) train data set
```

```
In [35]: round(churn_train_rebal_py['Churn'].sum()*100/churn_train_rebal_py.shape[0],2)
```

```
Out[35]: 19.99
```

## R

```

In [36]: %%R
# Identify the record indices we want to resample using which()
to.resample_r <- which(churn_train_r$Churn == 'True')

In [37]: %%R
# Randomly sample from the values in to.resample
our.resample_r <- sample(x = to.resample_r, size = 158, replace = TRUE)

In [38]: %%R
# Get the records whose record numbers are those in our.resample
our.resample_r <- churn_train_r[our.resample_r,]

In [39]: %%R
# Add the resampled records back onto our original training data set
churn_train_rebal_r <- rbind(churn_train_r, our.resample_r)

In [40]: %%R
t1 <- table(churn_train_rebal_r$Churn)
#ratio <- t1[2] / sum(t1) * 100
t1

```

```

False  True
1902    483

```

```

In [41]: %%R
# Now the new ratio of Churn == True
t1['True']*100/dim(churn_train_rebal_r)[1]

      True
20.25157

```

**32. Which baseline model do we use to compare our classification model performance against? To which value does this baseline model assign all predictions? What is the accuracy of this baseline model?**

Answer: For binary classification, we use All Positive Model or All Negative Model.

All Positive Model assigns all predictors to positive values; All Negative Model assigns all predictors to negative values.

The accuracy for All Positive Model will be p which is 14% from our test data set; The accuracy for All Negative Model will be 1-p which will be 86%.

**33. Validate your partition by testing for the difference in mean day minutes for the training set versus the test set.**

## Python



```
In [42]: # Before rebalance in train data set
mean_day_mins_train_before_py = churn_train_py['Day Mins'].sum() / churn_train_py
mean_day_mins_train_before_py
```

Out[42]: 179.75302283922974

```
In [43]: # Mean days min in test set
mean_day_mins_test_py = churn_test_py['Day Mins'].sum() / churn_test_py.shape[0]
mean_day_mins_test_py
```

Out[43]: 179.81990909090908

Means are very similar, so running t-test

```
In [44]: # Two samples t test
np.var(churn_test_py['Day Mins']), np.var(churn_train_py['Day Mins'])
```

Out[44]: (3010.977976355374, 2943.552898385955)

```
In [45]: # Perform two sample t-test with equal variances
stats.ttest_ind(a=churn_test_py['Day Mins'], b=churn_train_py['Day Mins'], equal_
```

Out[45]: Ttest\_indResult(statistic=0.03333175951101492, pvalue=0.9734120244442176)

p-value = 0.973 which is greater than 0.05, we fail to reject the null hypothesis of the test. We do not have sufficient evidence to say that the mean day minutes between the two data sets are different.

## R

```
In [46]: %%R
# Before rebalance in train data set
mean_day_mins_train_before_r <- sum(churn_train_r$Day.Mins) / dim(churn_train_r)[1]
mean_day_mins_train_before_r
```

[1] 180.911

```
In [47]: %%R
# After rebalance in train data set
mean_day_mins_train_rebal_r <- sum(churn_train_rebal_r$Day.Mins) / dim(churn_train_rebal_r)[1]
mean_day_mins_train_rebal_r
```

[1] 182.5936

```
In [48]: %%R
# Mean days min in test set
mean_day_mins_test_r <- sum(churn_test_r$Day.Mins) / dim(churn_test_r)[1]
mean_day_mins_test_r

[1] 177.4879
```

```
In [49]: %%R
# t-test
t.test(churn_test_r$Day.Mins, churn_train_r$Day.Mins)
```

#### Welch Two Sample t-test

```
data: churn_test_r$Day.Mins and churn_train_r$Day.Mins
t = -1.7248, df = 2259.7, p-value = 0.08471
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -7.315115  0.468881
sample estimates:
mean of x mean of y
 177.4879  180.9110
```

p-value = 0.085 which is greater than 0.05, we fail to reject the null hypothesis of the test. We do not have sufficient evidence to say that the mean day minutes between the two data sets are different.

### 34. Validate your partition by testing for the difference in proportion of true churn records for the training set versus the test set.

## Python

```
In [50]: # Ratio of "True" to all Churn in churn_train before rebalance
print(churn_train_py['Churn'].sum()*100 / churn_train_py.shape[0])

14.33049708911778
```

```
In [51]: # Ratio of "True" to all Churn in churn_test
print(churn_test_py['Churn'].sum()*100 / churn_test_py.shape[0])

14.818181818181818
```

Two proportions are similar, running one proportion z-test

```
In [52]: # Perform one proportion z-test
proportions_ztest(count=churn_train_py['Churn'].sum(), nobs=churn_train_py.shape[0])

Out[52]: (-0.6577180599814599, 0.5107193152727614)
```

p-value = 0.51 which is greater than 0.05, we fail to reject the null hypothesis of the test. We do not have sufficient evidence to say that the proportion of true churn records between the two data sets are different.

## R

```
In [53]: %%R
# Ratio of "True" to all Churn in churn_train before rebalance

table(churn_train_r$Churn)["True"]*100 / dim(churn_train_r)[1]

      True
14.59362
```

```
In [54]: %%R
# Ratio of "True" to all Churn in churn_test

table(churn_test_r$Churn)["True"]*100 / dim(churn_test_r)[1]

      True
14.28571
```

```
In [55]: %%R
# t-test
prop.test(c(table(churn_test_r$Churn)["True"],table(churn_train_r$Churn)["True"]))

      2-sample test for equality of proportions with continuity correction

data:  c(table(churn_test_r$Churn)["True"], table(churn_train_r$Churn)["True"])
out of c(dim(churn_test_r)[1], dim(churn_train_r)[1])
X-squared = 0.034423, df = 1, p-value = 0.8528
alternative hypothesis: two.sided
95 percent confidence interval:
 -0.02905977  0.02290158
sample estimates:
   prop 1    prop 2 
0.1428571 0.1459362
```

p-value = 0.85 which is greater than 0.05, we fail to reject the null hypothesis of the test. We do not have sufficient evidence to say that the proportion of true churn records between the two data sets are different.

**For the following exercises, work with the adult\_ch6\_training and adult\_ch6\_test data sets. Use R to solve each problem.**

## Python

```
In [56]: adult_train_py = pd.read_csv("D:/2021-Spring-textbooks/ADS-502/Website Data Sets/adult_train.csv")
adult_test_py = pd.read_csv("D:/2021-Spring-textbooks/ADS-502/Website Data Sets/adult_test.csv")

print(adult_train_py.head())
print(adult_test_py.head())
```

	Marital status	Income	Cap_Gains_Losses
0	Never-married	<=50K	0.02174
1	Divorced	<=50K	0.00000
2	Married	<=50K	0.00000
3	Married	<=50K	0.00000
4	Married	<=50K	0.00000

	Marital status	Income	Cap_Gains_Losses
0	Married	<=50K	0.000000
1	Married	>50K	0.051781
2	Never-married	<=50K	0.000000
3	Divorced	>50K	0.000000
4	Married	>50K	0.000000

## R

```
In [57]: %%R
adult_train_r <- read.csv(file = "D:/2021-Spring-textbooks/ADS-502/Website Data Sets/adult_train.csv")
adult_test_r <- read.csv(file = "D:/2021-Spring-textbooks/ADS-502/Website Data Sets/adult_test.csv")

print(head(adult_train_r))
print(head(adult_test_r))
```

	Marital.status	Income	Cap_Gains_Losses
1	Never-married	<=50K	0.02174
2	Divorced	<=50K	0.00000
3	Married	<=50K	0.00000
4	Married	<=50K	0.00000
5	Married	<=50K	0.00000
6	Married	>50K	0.00000

	Marital.status	Income	Cap_Gains_Losses
1	Married	<=50K	0.000000
2	Married	>50K	0.051781
3	Never-married	<=50K	0.000000
4	Divorced	>50K	0.000000
5	Married	>50K	0.000000
6	Married	<=50K	0.000000

```
In [58]: %%R
# Change categorical attributes to factors
adult_train_r$Income <- factor(adult_train_r$Income)
adult_train_r$Marital.status <- factor(adult_train_r$Marital.status)
```

**23. Using the training data set, create a C5.0 model (Model 1) to predict a customer's Income using Marital Status and Capital Gains and Losses. Obtain the predicted responses.**

## Python

```
In [59]: y = adult_train_py[['Income']]
y_test = adult_test_py[['Income']]
X = adult_train_py[['Marital status', 'Cap_Gains_Losses']]
X_test = adult_test_py[['Marital status', 'Cap_Gains_Losses']]
marital_dummy = pd.get_dummies(X[['Marital status']])
marital_dummy_test = pd.get_dummies(X_test[['Marital status']])
X = pd.concat((X[['Cap_Gains_Losses']], marital_dummy), axis = 1)
X_test = pd.concat((X_test[['Cap_Gains_Losses']], marital_dummy_test), axis = 1)
C5_py = DecisionTreeClassifier(criterion='entropy', min_samples_split=75, max_lea
C5_test_py = DecisionTreeClassifier(criterion='entropy', min_samples_split=75, ma
```

```
In [60]: # Select a random sample
sample_01 = X.sample()
sample_01
```

```
Out[60]:
```

	Cap_Gains_Losses	Divorced	Married	Never-married	Separated	Widowed
5059	0.383838	0	1	0	0	0

```
In [61]: # Prediction of the selected random sample
pred_single_C5_py = C5_py.predict(sample_01)
pred_single_C5_py
```

```
Out[61]: array(['>50K'], dtype=object)
```

```
In [62]: # Compare to its true value

y.iloc[sample_01.index[0]]
```

```
Out[62]: Income    <=50K
Name: 5059, dtype: object
```

## R

```
In [63]: %%R
# Run the training data set through C5.0 to obtain Model 1. Save the result as C5
C5_r <- C5.0(Income ~ Marital.status + Cap_Gains_Losses, data = adult_train_r)
```

```
In [64]: %%R
# Subset the predictor variables from the test data set into their own data frame
train.X <- subset(x = adult_train_r, select = c("Marital.status", "Cap_Gains_Losses"))

ypred <- predict(object = C5_r, newdata = train.X)

# Select a random sample for prediction
sample_r <- adult_train_r[sample(nrow(adult_train_r), size=1), ]
sample_r
```

```
Marital.status Income Cap_Gains_Losses
4431 Never-married <=50K 0
```

```
In [65]: %%R

# Run sample prediction
test.sample_r <- subset(x = sample_r, select = c("Marital.status", "Cap_Gains_Losses"))
ypred_sample <- predict(object = C5_r, newdata = test.sample_r)
ypred_sample
```

```
[1] <=50K
Levels: <=50K >50K
```

**24. Evaluate Model 1 using the test data set. Construct a contingency table to compare the actual and predicted values of Income.**

## Python

```
In [66]: y_true = adult_test_py[['Income']]
y_pred = C5_test_py.predict(X_test)
confusion_py = confusion_matrix(y_true, y_pred)
confusion_py
```

```
Out[66]: array([[4627,  47],
                [1141, 340]], dtype=int64)
```

## R

```
In [67]: %%R

test.X <- subset(x = adult_test_r, select = c("Marital.status", "Cap_Gains_Losses"))
ypred_test <- predict(object = C5_r, newdata = test.X)

# Build a contingency table and compare the actual Income in test data set with predicted

t1 <- table(adult_test_r$Income, ypred_test)
row.names(t1) <- c("Actual: 0", "Actual: 1")
colnames(t1) <- c("Predicted: 0", "Predicted: 1")
t1 <- addmargins(A = t1, FUN = list(Total = sum), quiet = TRUE)
t1
```

	ypred_test		
	Predicted: 0	Predicted: 1	Total
Actual: 0	4658	16	4674
Actual: 1	1057	424	1481
Total	5715	440	6155

**25. For Model 1, recapitulate Table 7.4 from the text, calculating all of the model evaluation measures shown in the table. Call this table the Model Evaluation Table. Leave space for Model 2.**

## Python

```
In [68]: TN_py = confusion_py[0,0]
FN_py = confusion_py[1,0]
FP_py = confusion_py[0,1]
TP_py = confusion_py[1,1]
TPN_py = TN_py + FN_py
TPP_py = FP_py + TP_py
TAN_py = TN_py + FP_py
TAP_py = FN_py + TP_py
GT_py = TN_py + FN_py + FP_py + TP_py
Precision_py = TP_py/TPP_py
Recall_py = TP_py/TAP_py
```

```
In [69]: # Overall model cost and Profit per customer formulas
Cost_FP = 10
Cost_TP = -40
Overall_model_cost = FP_py*Cost_FP + TP_py*Cost_TP
Profit_per_customer = -Overall_model_cost/GT_py
```

```
In [70]: column_names = ['Formula', 'Value']
row_names = ['Accuracy', 'Error rate', 'Sensitivity', 'Specificity', 'Precision',
              'Overall Model Cost', 'Profit Per Customer']
matrix = np.reshape(('(TN+TP)/GT', round((TN_py+TP_py)/GT_py,4), '1-(TN+TP)/GT', round(
              round(TP_py/TAP_py,4), 'TN/TAN', round(TN_py/TAN_py,4), 'TP/TP',
              '2*Precision*Recall/(Precision+Recall)', round(2*Precision_py
              '5*Precision*Recall/((4+Precision)+Recall)', round(5*Precision_py
              '1.25*Precision*Recall/((0.25*Precision)+Recall)',
              round(1.25*Precision_py*Recall_py/((0.25*Precision_py)+Recall_py),
              Overall_model_cost, '-Overall_model_cost/GT', round(Profit_per_
Model_Evaluation_Table_py = pd.DataFrame(matrix, columns=column_names, index=row_
Model_Evaluation_Table_py
```

Out[70]:

	Formula	Value
<b>Accuracy</b>	$(TN+TP)/GT$	0.807
<b>Error rate</b>	$1-(TN+TP)/GT$	0.193
<b>Sensitivity</b>	$TP/TAP$	0.2296
<b>Specificity</b>	$TN/TAN$	0.9899
<b>Precision</b>	$TP/TPP$	0.8786
<b>F1</b>	$2*Precision*Recall/(Precision+Recall)$	0.364
<b>F2</b>	$5*Precision*Recall/((4+Precision)+Recall)$	0.1974
<b>F0.5</b>	$1.25*Precision*Recall/((0.25*Precision)+Recall)$	0.5612
<b>Overall Model Cost</b>	$FP*Cost\_FP + TP*Cost\_TP$	-13130
<b>Profit Per Customer</b>	$-Overall\_model\_cost/GT$	2.1332

## R

```
In [71]: %%R
# Define elements in the Model Evaluation Table
TN_r <- t1[1,1]
FP_r <- t1[1,2]
FN_r <- t1[2,1]
TP_r <- t1[2,2]
TPN_r <- t1[3,1]
TPP_r <- t1[3,2]
TAN_r <- t1[1,3]
TAP_r <- t1[2,3]
GT_r <- t1[3,3]
Precision_r <- TP_r/TPP_r
Recall_r <- TP_r/TAP_r
```



```
In [72]: %%R
# Overall model cost and Profit per customer formulas
Cost_FP = 10
Cost_TP = -40
Overall_model_cost_r = FP_r*Cost_FP + TP_r*Cost_TP
Profit_per_customer_r = -Overall_model_cost_r/GT_r
```

```
In [73]: %%R
tab <- matrix(c('(TN+TP)/GT', round((TN_r+TP_r)/GT_r,4), '1-(TN+TP)/GT', round(1-(TN_r+TP_r)/GT_r,4),
'TN/TAN', round(TN_r/TAN_r,4), 'TP/TPP', round(TP_r/TPP_r,4), '2*Precision*Recall/(Precision+Recall)',
round(2*Precision_r*Recall_r/(Precision_r+Recall_r),4), '5*Precision*Recall/((4+Precision)+Recall)',
round(5*Precision_r*Recall_r/((4+Precision_r)+Recall_r),4), '1.25*Precision*Recall/((0.25*Precision)+Recall)',
round(1.25*Precision_r*Recall_r/((0.25*Precision_r)+Recall_r),4),
Overall_model_cost_r, '-Overall_model_cost/GT', round(Profit_per_customer_r,4)),
colnames(tab) <- c('Formula', 'Value')
rownames(tab) <- c('Accuracy', 'Error rate', 'Sensitivity', 'Specificity', 'Precision', 'F1', 'F2', 'F0.5',
'Profit Per Customer')
tab <- as.table(tab)
Model_Evaluation_Table_r <- tab
Model_Evaluation_Table_r
```

	Formula	Value
Accuracy	(TN+TP)/GT	0.8257
Error rate	1-(TN+TP)/GT	0.1743
Sensitivity	TP/TAP	0.2863
Specificity	TN/TAN	0.9966
Precision	TP/TPP	0.9636
F1	2*Precision*Recall/(Precision+Recall)	0.4414
F2	5*Precision*Recall/((4+Precision)+Recall)	0.2627
F0.5	1.25*Precision*Recall/((0.25*Precision)+Recall)	0.6541
Overall Model Cost	FP*Cost_FP + TP*Cost_TP	-16800
Profit Per Customer	-Overall_model_cost/GT	2.7295

## 26. Clearly and completely interpret each of the Model 1 evaluation measures from the Model Evaluation Table.

Accuracy (ACC) is calculated as the number of all correct predictions divided by the total number of the dataset. The best accuracy is 1.0, whereas the worst is 0.0.

Error rate (ERR) is calculated as the number of all incorrect predictions divided by the total number of the dataset. The best error rate is 0.0, whereas the worst is 1.0.

Sensitivity (Recall or True positive rate) is calculated as the number of correct positive predictions divided by the total number of positives. It is also called recall (REC) or true positive rate (TPR). The best sensitivity is 1.0, whereas the worst is 0.0.

Specificity (True negative rate) is calculated as the number of correct negative predictions divided by the total number of negatives. It is also called true negative rate (TNR). The best specificity is 1.0, whereas the worst is 0.0.

Precision (Positive predictive value) is calculated as the number of correct positive predictions divided by the total number of positive predictions. It is also called positive predictive value (PPV). The best precision is 1.0, whereas the worst is 0.0.

F1-score is the harmonic mean of the precision and recall. A perfect model has an F-score of 1. A factor in  $F\beta$  indicating how much more important recall is than precision. For example, if we consider recall to be twice as important as precision, we can set  $\beta$  to 2. The standard F-score is equivalent to setting  $\beta$  to 1.

**27. Create a cost matrix, called the 3x cost matrix, that specifies a false positive is four times as bad as a false negative.**

## Python

```
In [74]: # Cost matrix with FP = 4FN
cost_matrix_3x_py = {'<=50K':1, '>50K':4}
```

## R

```
In [75]: %%R
cost_matrix_3x_r = matrix(c(0,4,1,0), byrow = TRUE, ncol = 2 )
dimnames(cost_matrix_3x_r) <- list(c("<=50K", ">50K"), c("<=50K", ">50K"))
cost_matrix_3x_r
```

```
      <=50K >50K
<=50K      0      4
>50K       1      0
```

**28. Using the training data set, build a C5.0 model (Model 2) to predict a customer's Income using Marital Status and Capital Gains and Losses, using the 3x cost matrix.**

## Python

```
In [76]: C5_02_py = DecisionTreeClassifier(criterion='entropy', min_samples_split=75, max_
                                             class_weight=cost_matrix_3x_py).fit(X, y)

# Select a random customer from the training set
sample_02 = X.sample()
sample_02
```

Out[76]:

	Cap_Gains_Losses	Divorced	Married	Never-married	Separated	Widowed
18201	0.0	1	0	0	0	0

```
In [77]: # Select a random customer from the training set
sample_02 = X.sample()
sample_02
```

```
Out[77]:
```

	Cap_Gains_Losses	Divorced	Married	Never-married	Separated	Widowed
12042	0.0	1	0	0	0	0

```
In [78]: # Predict Income for the single random customer
pred_single_C5_02_py = C5_02_py.predict(sample_02)
pred_single_C5_02_py
```

```
Out[78]: array(['<=50K'], dtype=object)
```

```
In [79]: # Compare to the true value

y.iloc[sample_02.index[0]]
```

```
Out[79]: Income    <=50K
Name: 12042, dtype: object
```

## R

```
In [80]: %%R
#Adding cost matrix to C5.0 model
C5.costs_r <- C5.0(Income ~ Marital.status + Cap_Gains_Losses, data = adult_train_r)
```

```
In [81]: %%R
# Select a random sample for prediction
sample_02_r <- adult_train_r[sample(nrow(adult_train_r), size=1), ]
sample_02_r
```

```
      Marital.status Income Cap_Gains_Losses
10104      Married  <=50K                0
```

```
In [82]: %%R
# Run sample prediction
test.sample_r <- subset(x = sample_02_r, select = c("Marital.status", "Cap_Gains_Losses"))
ypred_sample <- predict(object = C5.costs_r, newdata = test.sample_r)
ypred_sample
```

```
[1] >50K
Levels: <=50K >50K
```

**29. Evaluate your predictions from Model 2 using the actual response values from the test data set. Add Overall Model Cost and Profit per Customer to the Model Evaluation Table. Calculate all the measures from the Model Evaluation Table.**

## Python

```
In [83]: # Contingency table with Model 2 using test data
y_true = adult_test_py[['Income']]
y_pred_02 = C5_02_py.predict(X_test)
confusion_02_py = confusion_matrix(y_true, y_pred_02)
confusion_02_py
```

```
Out[83]: array([[3041, 1633],
               [ 149, 1332]], dtype=int64)
```

```
In [84]: TN_02_py = confusion_02_py[0,0]
FN_02_py = confusion_02_py[1,0]
FP_02_py = confusion_02_py[0,1]
TP_02_py = confusion_02_py[1,1]
TPN_02_py = TN_02_py + FN_02_py
TPP_02_py = FP_02_py + TP_02_py
TAN_02_py = TN_02_py + FP_02_py
TAP_02_py = FN_02_py + TP_02_py
GT_02_py = TN_02_py + FN_02_py + FP_02_py + TP_02_py
Precision_02_py = TP_02_py/TPP_02_py
Recall_02_py = TP_02_py/TAP_02_py
```

```
In [85]: # Overall model cost and Profit per customer formulas
Cost_FP = 10
Cost_TP = -40
Overall_model_cost_02 = FP_02_py*Cost_FP + TP_02_py*Cost_TP
Profit_per_customer_02 = -Overall_model_cost_02/GT_02_py
```

```
In [86]: # Model Evaluation Table with Overall model cost and Profit per customer
column_names = ['Formula', 'Value']
row_names     = ['Accuracy', 'Error rate', 'Sensitivity', 'Specificity', 'Precision',
                 'Overall Model Cost', 'Profit Per Customer']
matrix = np.reshape((('TN+TP)/GT', round((TN_02_py+TP_02_py)/GT_02_py,4), '1-(TN+TP)/GT',
                    round(TP_02_py/TAP_02_py,4), 'TN/TAN', round(TN_02_py/TAN_02_py,4),
                    '2*Precision*Recall/(Precision+Recall)', round(2*Precision_02_py*Recall_02_py/(Precision_02_py+Recall_02_py),4),
                    '5*Precision*Recall/((4+Precision)+Recall)', round(5*Precision_02_py*Recall_02_py/((4+Precision_02_py)+Recall_02_py),4),
                    '1.25*Precision*Recall/((0.25*Precision)+Recall)', round(1.25*Precision_02_py*Recall_02_py/((0.25*Precision_02_py)+Recall_02_py),4),
                    Overall_model_cost_02, '-Overall_model_cost/GT', round(Profit_Per_Customer_02/GT_02_py,4)),
                    (len(column_names), len(row_names)))
Model_Evaluation_Table_02_py = pd.DataFrame(matrix, columns=column_names, index=row_names)
Model_Evaluation_Table_02_py
```

Out[86]:

	Formula	Value
<b>Accuracy</b>	(TN+TP)/GT	0.7105
<b>Error rate</b>	1-(TN+TP)/GT	0.2895
<b>Sensitivity</b>	TP/TAP	0.8994
<b>Specificity</b>	TN/TAN	0.6506
<b>Precision</b>	TP/TPP	0.4492
<b>F1</b>	2*Precision*Recall/(Precision+Recall)	0.5992
<b>F2</b>	5*Precision*Recall/((4+Precision)+Recall)	0.3777
<b>F0.5</b>	1.25*Precision*Recall/((0.25*Precision)+Recall)	0.4992
<b>Overall Model Cost</b>	FP*Cost_FP + TP*Cost_TP	-36950
<b>Profit Per Customer</b>	-Overall_model_cost/GT	6.0032

## R

```
In [87]: %%R
ypred_test <- predict(object = C5.costs_r, newdata = test.X)

# Build a contingency table and compare the actual Income in test data set with predicted Income

t2 <- table(adult_test_r$Income, ypred_test)
row.names(t2) <- c("Actual: 0", "Actual: 1")
colnames(t2) <- c("Predicted: 0", "Predicted: 1")
t2 <- addmargins(A = t2, FUN = list(Total = sum), quiet = TRUE)
t2
```

	ypred_test		
	Predicted: 0	Predicted: 1	Total
Actual: 0	3171	1503	4674
Actual: 1	128	1353	1481
Total	3299	2856	6155

In [88]: %%R

# Define elements in the Model Evaluation Table 02

```

TN_02_r <- t2[1,1]
FP_02_r <- t2[1,2]
FN_02_r <- t2[2,1]
TP_02_r <- t2[2,2]
TPN_02_r <- t2[3,1]
TPP_02_r <- t2[3,2]
TAN_02_r <- t2[1,3]
TAP_02_r <- t2[2,3]
GT_02_r <- t2[3,3]
Precision_02_r <- TP_02_r/TPP_02_r
Recall_02_r <- TP_02_r/TAP_02_r

```

In [89]: %%R

# Overall model cost and Profit per customer formulas

```

Cost_FP = 10
Cost_TP = -40
Overall_model_cost_02_r = FP_02_r*Cost_FP + TP_02_r*Cost_TP
Profit_per_customer_02_r = -Overall_model_cost_02_r/GT_02_r

```

In [90]: %%R

```

tab2 <- matrix(c('(TN+TP)/GT', round((TN_02_r+TP_02_r)/GT_02_r,4), '1-(TN+TP)/GT',
                'TP/TAP', round(TP_02_r/TAP_02_r,4),
                'TN/TAN', round(TN_02_r/TAN_02_r,4), 'TP/TPP', round(TP_02_r/TPP_02_r,4),
                round(2*Precision_02_r*Recall_02_r/(Precision_02_r+Recall_02_r),4),
                round(5*Precision_02_r*Recall_02_r/((4+Precision_02_r)+Recall_02_r),4),
                round(1.25*Precision_02_r*Recall_02_r/((0.25*Precision_02_r)+Recall_02_r),4),
                Overall_model_cost_02_r, '-Overall_model_cost/GT', round(Profit_per_customer_02_r,4)),
              nrow=9, byrow=TRUE)
colnames(tab2) <- c('Formula', 'Value')
rownames(tab2) <- c('Accuracy', 'Error rate', 'Sensitivity', 'Specificity', 'Precision',
                  'F1', 'F2', 'F0.5', 'Profit Per Customer')
tab2 <- as.table(tab2)
Model_Evaluation_Table_02_r <- tab2
Model_Evaluation_Table_02_r

```

	Formula	Value
Accuracy	(TN+TP)/GT	0.735
Error rate	1-(TN+TP)/GT	0.265
Sensitivity	TP/TAP	0.9136
Specificity	TN/TAN	0.6784
Precision	TP/TPP	0.4737
F1	2*Precision*Recall/(Precision+Recall)	0.6239
F2	5*Precision*Recall/((4+Precision)+Recall)	0.4017
F0.5	1.25*Precision*Recall/((0.25*Precision)+Recall)	0.5242
Overall Model Cost	FP*Cost_FP + TP*Cost_TP	-39090
Profit Per Customer	-Overall_model_cost/GT	6.3509

**30. Compare the evaluation measures from Model 1 and Model 2 using the 3x cost matrix. Discuss the strengths and weaknesses of each model.**

## Python

```
In [91]: Model_Comparison_py = Model_Evaluation_Table_py.drop(['Formula'], axis=1)
```

```
In [92]: Model_Comparison_py['Value of Model 2'] = Model_Evaluation_Table_02_py[['Value']]
```

```
In [93]: Model_Comparison_py
```

Out[93]:

	Value	Value of Model 2
<b>Accuracy</b>	0.807	0.7105
<b>Error rate</b>	0.193	0.2895
<b>Sensitivity</b>	0.2296	0.8994
<b>Specificity</b>	0.9899	0.6506
<b>Precision</b>	0.8786	0.4492
<b>F1</b>	0.364	0.5992
<b>F2</b>	0.1974	0.3777
<b>F0.5</b>	0.5612	0.4992
<b>Overall Model Cost</b>	-13130	-36950
<b>Profit Per Customer</b>	2.1332	6.0032

Compare Model 1 to Model 2 we can see that Model 1 has higher accuracy and lower error rate while Model 2 scores higher on F-scores. Adding the unequal error costs makes Model 2 has decreased overall model cost and increased the Profit Per Customer by 181%. Because every true positive gives us \$40, the models which tended to make more positive predictions did better. Sensitivity (recall), the proportion of all the positive responders that the model captured, thus turned out to be more important than specificity. Thus, accuracy is not the proper metric to compare models which have unequal error costs.

## R

```
In [94]: %%R
Model_Evaluation_Table_r <- Model_Evaluation_Table_r[, -1]
```

```
In [95]: %%R
Model_Comparison_r <- cbind(Model_Evaluation_Table_r, Model_Evaluation_Table_02_r)
```

```
In [96]: %%R
Model_Comparison_r
```

```

                                Model_Evaluation_Table_r
Accuracy                        "0.8257"                "0.735"
Error rate                     "0.1743"                "0.265"
Sensitivity                    "0.2863"                "0.9136"
Specificity                   "0.9966"                "0.6784"
Precision                     "0.9636"                "0.4737"
F1                            "0.4414"                "0.6239"
F2                            "0.2627"                "0.4017"
F0.5                         "0.6541"                "0.5242"
Overall Model Cost            "-16800"                 "-39090"
Profit Per Customer           "2.7295"                 "6.3509"
```

Same scenario with Python.

**For the following exercises, work with the `framingham_nb_training` and `framingham_nb_test` data sets. Use either Python or R to solve each problem.**

**31. Run the Naïve Bayes classifier to classify persons as living or dead based on sex and education.**

## Python

```
In [97]: framingham_nb_train_py = pd.read_csv("D:/2021-Spring-textbooks/ADS-502/Website Data/framingham_nb_train.csv")
framingham_nb_test_py = pd.read_csv("D:/2021-Spring-textbooks/ADS-502/Website Data/framingham_nb_test.csv")

framingham_nb_train_py.head()
```

Out[97]:

	Sex	Educ	Death
0	2	3	0
1	2	2	0
2	1	1	0
3	2	1	0
4	2	1	0

```
In [98]: # Check for if the true response has a good proportion
framingham_nb_train_py['Death'].sum()*100 / framingham_nb_train_py.shape[0]
```

Out[98]: 55.0

True response rate is 55%, no need to rebalance the data. Move on.



```
In [99]: # Split the predictors and target in the training set
X_framingham_train_py = framingham_nb_train_py[['Sex', 'Educ']]
y_framingham_train_py = framingham_nb_train_py['Death']
```

```
In [100]: # Convert categorical data into dummy variables for both 'Sex' and 'Educ' for the
X_framingham_train_processed_py = pd.DataFrame()

for var in X_framingham_train_py.columns:
    dummies = pd.get_dummies(X_framingham_train_py[var])
    X_framingham_train_processed_py = pd.concat([X_framingham_train_processed_py,
```

```
In [101]: # Run Naive Bayes algorithm using training set
nb = MultinomialNB().fit(X_framingham_train_processed_py, y_framingham_train_py)
```

```
In [102]: # Split the predictors and target in the test set
X_framingham_test_py = framingham_nb_test_py[['Sex', 'Educ']]
y_framingham_test_py = framingham_nb_test_py['Death']
```

```
In [103]: # Convert categorical data into dummy variables for both 'Sex' and 'Educ' for the
X_framingham_test_processed_py = pd.DataFrame()

for var in X_framingham_test_py.columns:
    dummies = pd.get_dummies(X_framingham_test_py[var])
    X_framingham_test_processed_py = pd.concat([X_framingham_test_processed_py, c
```

```
In [104]: predictions_nb_py = nb.predict(X_framingham_test_processed_py)
          predictions_nb_py
```

```
Out[104]: array([1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0,
                1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1,
                1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0,
                0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1,
                1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0,
                1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1,
                0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1,
                1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1,
                1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1,
                1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0,
                0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0,
                1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1,
                1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0,
                1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1,
                0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0,
                0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0,
                1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1,
                1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1,
                1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1,
                1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1,
                1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1,
                1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
                1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
                1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
                1, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1,
                0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1,
                0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
                1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0], dtype=int64)
```

**R**

In [105]: %%R

```
framingham_nb_train_r <- read.csv(file = "D:/2021-Spring-textbooks/ADS-502/Website  
framingham_nb_test_r <- read.csv(file = "D:/2021-Spring-textbooks/ADS-502/Website  
  
head(framingham_nb_train_r)
```

	Sex	Educ	Death
1	2	3	0
2	2	2	0
3	1	1	0
4	2	1	0
5	2	1	0
6	2	3	0

In [106]: %%R

```
cols = c('Sex', 'Educ')  
framingham_nb_train_r[, cols] <- lapply(framingham_nb_train_r[, cols], as.factor)  
  
# Or we could use this one by one  
#framingham_nb_train_r$Sex <- factor(framingham_nb_train_r$Sex)  
#framingham_nb_train_r$Educ <- factor(framingham_nb_train_r$Educ)
```

In [107]: %%R

```
# Run the model  
  
nb <- naiveBayes(formula = Death ~ Sex + Educ, data = framingham_nb_train_r)
```

In [108]: %%R

```
# Process the test set the same way as training set  
framingham_nb_test_r[, cols] <- lapply(framingham_nb_test_r[, cols], as.factor)
```

```
In [109]: %%R
predictions_nb_r <- predict(nb, framingham_nb_test_r)
predictions_nb_r
```

```
[1] 0 0 1 1 1 0 1 1 1 1 0 1 1 0 0 0 0 1 0 1 0 0 1 1 1 1 0 0 1 1 1 0 1 0
0
[38] 1 1 1 0 1 1 0 0 0 1 1 1 0 1 1 1 0 1 1 1 0 1 0 1 1 0 1 1 0 0 1 1 1 1
1
[75] 1 0 0 0 0 0 1 0 1 0 1 0 1 1 1 1 1 1 1 0 1 1 0 0 0 1 1 1 1 1 1 1 0 1 1
0
[112] 1 1 1 1 1 0 1 1 1 1 1 0 1 0 0 1 1 1 0 1 1 1 1 0 0 1 0 1 1 0 0 1 1 1 0 1
1
[149] 0 0 0 1 1 1 0 0 0 0 1 0 1 0 0 1 1 1 0 1 1 0 1 1 1 0 1 1 0 1 0 0 1 0 1 1
1
[186] 0 0 0 1 0 1 0 1 1 1 0 1 1 1 0 0 1 1 1 0 1 1 1 0 0 1 1 0 1 1 0 0 1 0 1 1
0
[223] 0 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 0 1 1 0 0 0 0 1 1 1 1 1 0 1 0 0 0 1 1 1 0
0
[260] 0 1 1 1 1 1 0 1 1 1 0 1 1 1 1 1 1 0 1 0 0 1 1 1 0 1 1 0 1 1 0 1 1 1 0 1
1
[297] 1 1 1 0 1 1 1 1 1 1 1 0 1 1 1 0 0 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0
0
[334] 0 0 0 1 1 1 1 0 0 1 0 1 1 0 0 1 0 0 1 1 0 1 0 1 0 0 0 1 0 0 0 1 1 0 0 1
0
[371] 1 1 0 1 1 1 1 0 1 0 1 0 0 1 1 1 1 0 1 1 1 1 1 0 0 0 1 1 0 0 0 1 1 0 1 0
1
[408] 1 0 0 0 0 1 0 0 1 1 1 0 0 1 0 1 1 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 1 1 1 1
1
[445] 1 1 1 1 1 0 1 1 0 1 1 1 1 1 1 1 0 1 1 0 1 1 0 1 0 1 1 1 1 1 1 1 1 0 1 0
1
[482] 1 1 1 1 1 0 1 1 1 1 1 1 0 0 1 1 1 0 1 0 0 1 0 1 1 1 0 1 1 0 1 1 1 0 1 1
1
[519] 1 1 0 0 0 1 0 1 1 0 0 1 1 1 1 1 1 0 0 1 1 1 1 0 0 0 1 0 0 1 1 1 1 1 1 0
1
[556] 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0 1 1 0 1 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 0
1
[593] 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 1 0 1 1 0 0 0 1 1 1 1 1 0 1 1 0
1
[630] 1 1 1 1 0 1 0 1 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1
1
[667] 1 1 0 0 0 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 0 1 1 1 0 1
1
[704] 1 1 1 1 1 1 0 1 0 1 1 0 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0 1
1
[741] 1 1 1 0 1 1 1 0 1 1 1 1 1 0 1 1 1 1 0 1 1 1 1 1 0 1 0 1 0 1 0 1 1 1 1 1
1
[778] 1 1 1 1 1 1 1 0 1 0 0 0 1 1 1 1 1 1 0 0 1 1 0 1 1 0 1 1 1 0 1 1 1 1 0 0
1
[815] 0 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 0 1 1 1 1 0 0 1 1 1 1 1 1
1
[852] 1 1 0 1 1 1 0 1 1 0 1 1 0 1 1 1 1 0 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1
1
[889] 1 1 1 1 1 1 0 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 0 0 1 1 1
1
[926] 1 1 1 1 1 0 0 1 1 1 1 1 0 1 1 1 1 1 1 0 1 1 0 1 1 1 1 1 1 1 1 0 1 0 1 1 1
1
[963] 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1
```

```
0
[1000] 0
Levels: 0 1
```

**32. Evaluate the Naïve Bayes model on the `framingham_nb_test` data set. Display the results in a contingency table. Edit the row and column names of the table to make the table more readable. Include a total row and column.**

## Python

```
In [110]: cm = confusion_matrix(y_framingham_test_py, predictions_nb_py)
```

```
In [111]: cm
```

```
Out[111]: array([[203, 322],
                 [105, 370]], dtype=int64)
```

```
In [112]: TN_py = cm[0][0]
FP_py = cm[0][1]
FN_py = cm[1][0]
TP_py = cm[1][1]

print('TN: ', TN_py,
      '\nFP: ', FP_py,
      '\nFN: ', FN_py,
      '\nTP: ', TP_py)
```

```
TN: 203
FP: 322
FN: 105
TP: 370
```

```
In [113]: col_names = ['Predicted:0', 'Predicted:1', 'Total']
row_names = ['Actual:0', 'Actual:1', 'Total']
matrix = np.reshape((TN_py, FP_py, TN_py+FP_py, FN_py, TP_py, FN_py+TP_py, TN_py+FN_py),
cm_py = pd.DataFrame(matrix, columns=col_names, index=row_names)
cm_py
```

```
Out[113]:
```

	Predicted:0	Predicted:1	Total
Actual:0	203	322	525
Actual:1	105	370	475
Total	308	692	1000

## R

```
In [114]: %%R
# Build a contingency table and compare the actual Income in test data set with p

t3 <- table(framingham_nb_test_r$Death, predictions_nb_r)
row.names(t3) <- c("Actual: 0", "Actual: 1")
colnames(t3) <- c("Predicted: 0", "Predicted: 1")
t3 <- addmargins(A = t3, FUN = list(Total = sum), quiet = TRUE)
```

```
In [115]: %%R
t3
```

	predictions_nb_r		
	Predicted: 0	Predicted: 1	Total
Actual: 0	167	283	450
Actual: 1	122	428	550
Total	289	711	1000

```
In [116]: %%R
TN_r = t3[1:1,1:1]
FP_r = t3[1:1,2:2]
FN_r = t3[2:2,1:1]
TP_r = t3[2:2,2:2]
```

```
In [117]: %%R
cat(paste("TN:", FP_r, "\nFP:", FP_r, "\nFN:", FN_r, "\nTP:", TP_r))
```

```
TN: 283
FP: 283
FN: 122
TP: 428
```

**33. According to your table in the previous exercise, find the following values for the Naïve Bayes model:**

**a. Accuracy**

**b. Error rate**

## Python

```
In [118]: print("Accuracy for NB model is: ", round((TN_py+TP_py)/(TN_py+TP_py+FN_py+FP_py),
'\nError Rate for NB model is: ', round(1-(TN_py+TP_py)/(TN_py+TP_py+FN_py+FP_py),
```

```
Accuracy for NB model is: 0.573
Error Rate for NB model is: 0.427
```

## R

```
In [119]: %%R
cat(paste("Accuracy for NB model is: ", round((TN_r+TP_r)/(TN_r+TP_r+FN_r+FP_r),4),
"\nError Rate for NB model is: ", round(1-(TN_r+TP_r)/(TN_r+TP_r+FN_r+FP_r),4),
"\n"))

Accuracy for NB model is:  0.595
Error Rate for NB model is:  0.405
```

**34. According to your contingency table, find the following values for the Naïve Bayes model:**

**a. How often it correctly classifies dead persons.**

**b. How often it correctly classifies living persons.**

## Python

```
In [120]: print(round(TN_py*100/(TN_py+FP_py),2), "% chance of correctly classifies dead per
38.67 % chance of correctly classifies dead persons. This is also the Specifici
ty
```

```
In [121]: print(round(TP_py*100/(TP_py+FN_py),2), "% chance of correctly classifies living p
77.89 % chance of correctly classifies living persons. This is also the Sensiti
vity
```

## R

```
In [122]: %%R
cat(round(TN_r*100/(TN_r+FP_r),2), "% chance of correctly classifies dead persons.
37.11 % chance of correctly classifies dead persons. This is also the Specifici
ty
```

```
In [123]: %%R
cat(round(TP_r*100/(TP_r+FN_r),2), "% chance of correctly classifies living person
77.82 % chance of correctly classifies living persons. This is also the Sensiti
vity
```